

Introduction to Genomic Selection in R using the rrBLUP Package

Amy Jacobson
jaco0795@umn.edu
University of Minnesota



Learning Objectives

- Download the package and load the sample files
- Impute missing markers using `A.mat()`
- Define the training and validation populations
- Run `mixed.solve()` and determine accuracy of predictions

Overview of rrBLUP package

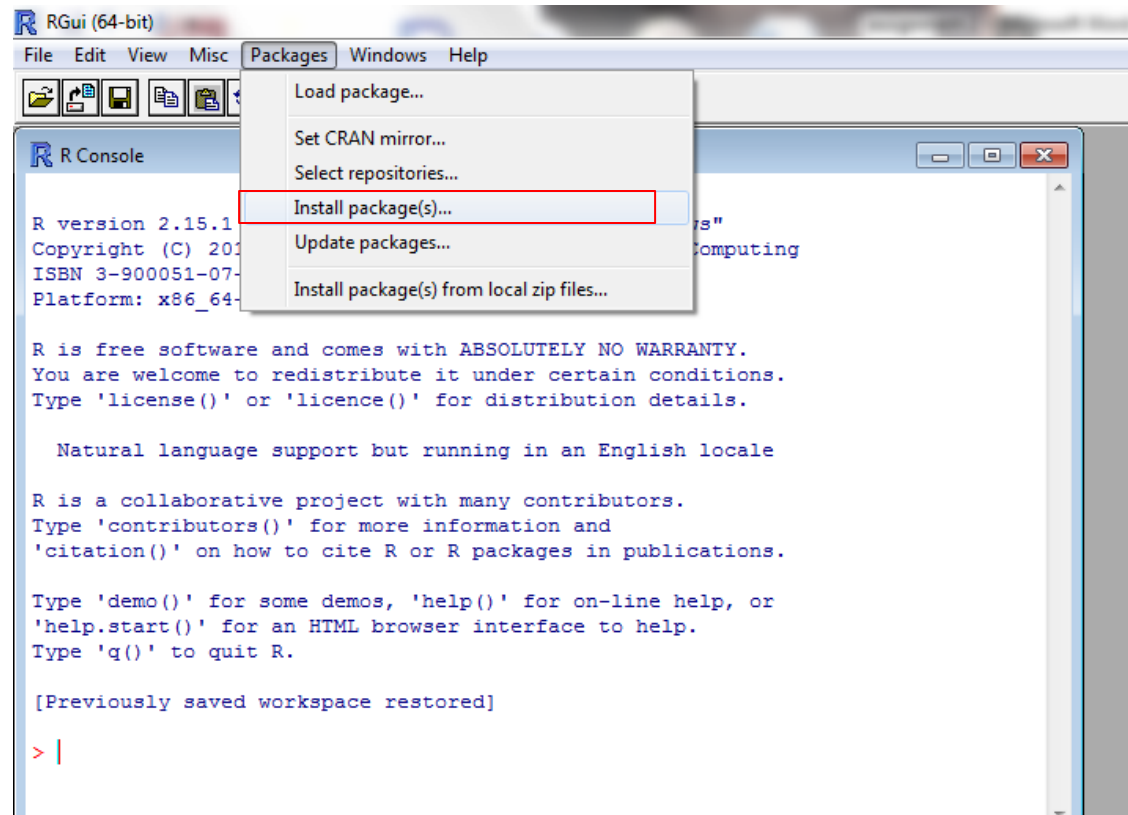
- Download from CRAN-version 4
 - Must use R version 2.14.1 or greater
- Uses ridge regression BLUP for genomic predictions
- Predicts marker effects through `mixed.solve()`
- `A.mat()` command can be used to impute missing markers
 - `Mixed.solve` does not allow NA marker values
- Define the training and validation populations

One Step vs. Two Step

- One step
 - Uses a mixed model analysis for the plot data
- Two step
 - Adjusted means are calculated across locations
 - Means are then used in ridge regression blup
- This webinar uses a two step approach
 - Computationally more efficient and faster

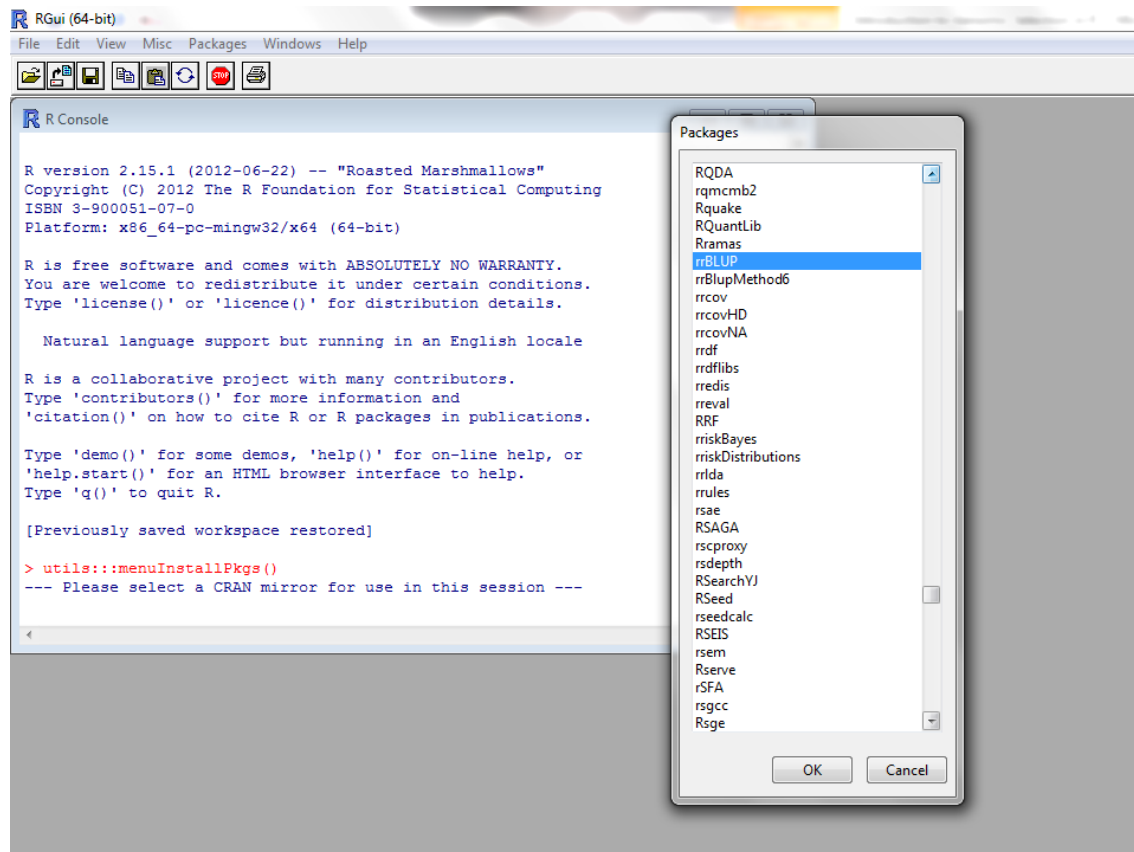
Install the rrBLUP Package

- Launch R->Packages->Install Package
- Select CRAN Mirror nearest you



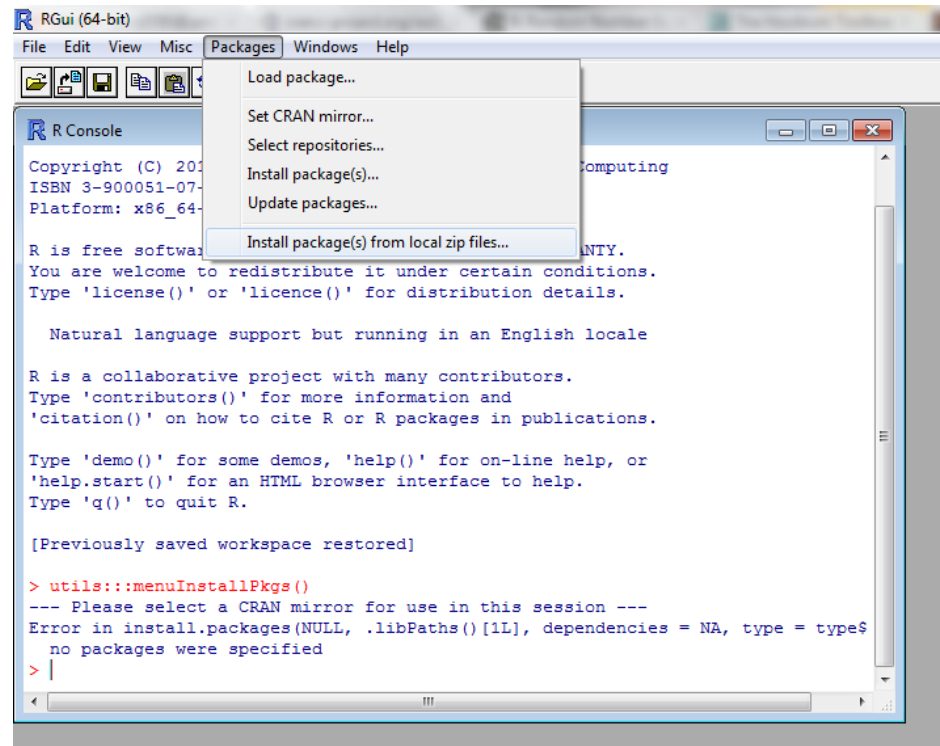
Install the rrBLUP Package

- Select the rrBLUP package




Install the rrBLUP Package

- Install the package by a zip file
- <http://cran.r-project.org/web/packages/rrBLUP/index.html>
- Packages->install package from local zip files



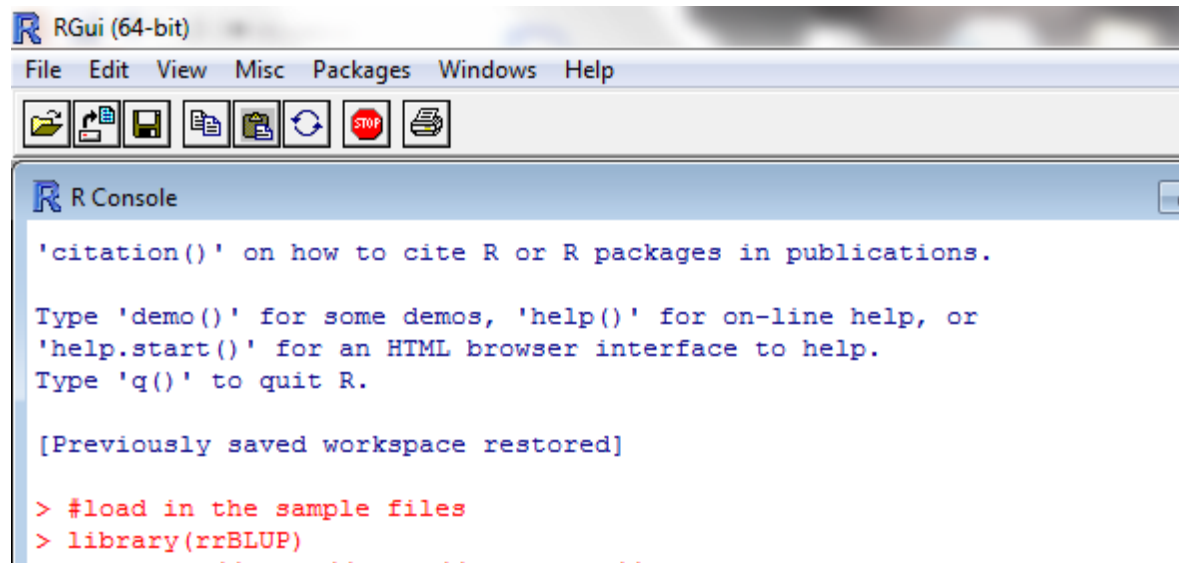
Install the rrBLUP Package

- Select the package from saved location



Install the rrBLUP Package


- Now that the package is installed, the library must be loaded every time R is opened



The screenshot shows the RGui (64-bit) window. The menu bar includes File, Edit, View, Misc, Packages, Windows, and Help. Below the menu bar is a toolbar with icons for file operations and execution. The R Console window is open, displaying the following text:

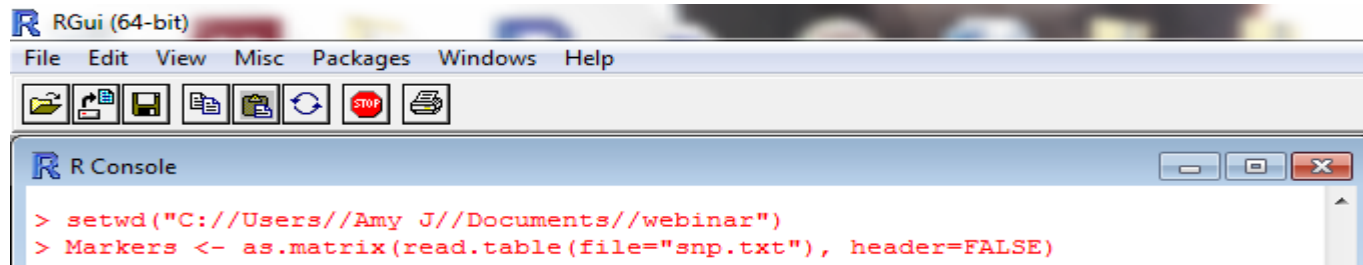
```
'citation()' on how to cite R or R packages in publications.  
Type 'demo()' for some demos, 'help()' for on-line help, or  
'help.start()' for an HTML browser interface to help.  
Type 'q()' to quit R.  
  
[Previously saved workspace restored]  
  
> #load in the sample files  
> library(rrBLUP)  
... ..
```

Sample Files

- Files downloaded from the Hordeum Toolbox
<http://hordeumtoolbox.org/>
- University of Minnesota barley breeding program preliminary yield trial-St. Paul location in 2009
- Phenotypic traits-yield, plant height and heading date
- 1178 markers, 164 NA markers
- 1 = homozygous for parent 1, 0 = heterozygous, and -1 homozygous for parent 2
 - Markers must be in the {-1,0,1} format for rrBLUP

Load the Sample Files

- Setwd()-Set the working directory to the location of the sample files
- Read.table command used for .txt files
- Read.csv command used for .csv files
- Header=F since sample marker file does not have a header with marker names

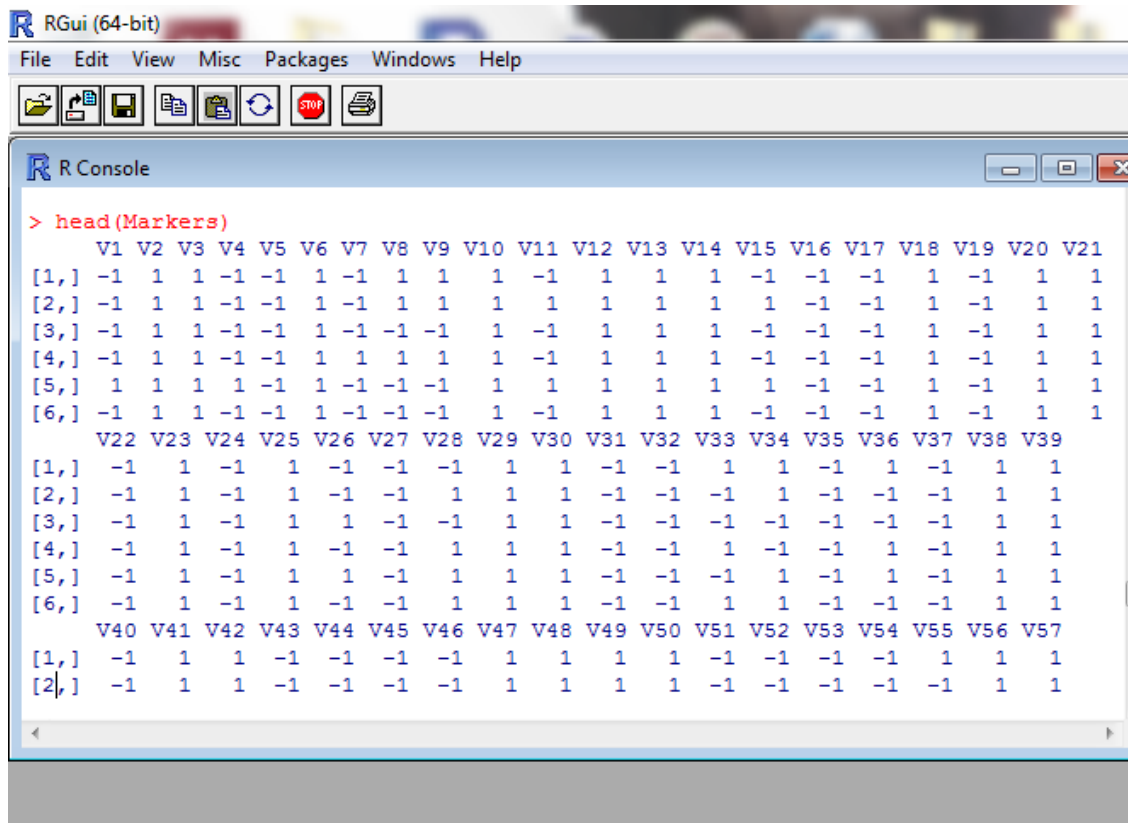


```
RGui (64-bit)
File Edit View Misc Packages Windows Help

R Console
> setwd("C://Users//Amy J//Documents//webinar")
> Markers <- as.matrix(read.table(file="snp.txt"), header=FALSE)
```

Load the Sample Files

- head() command used to see the first 5 lines of a file
- Useful to see if data was loaded correctly

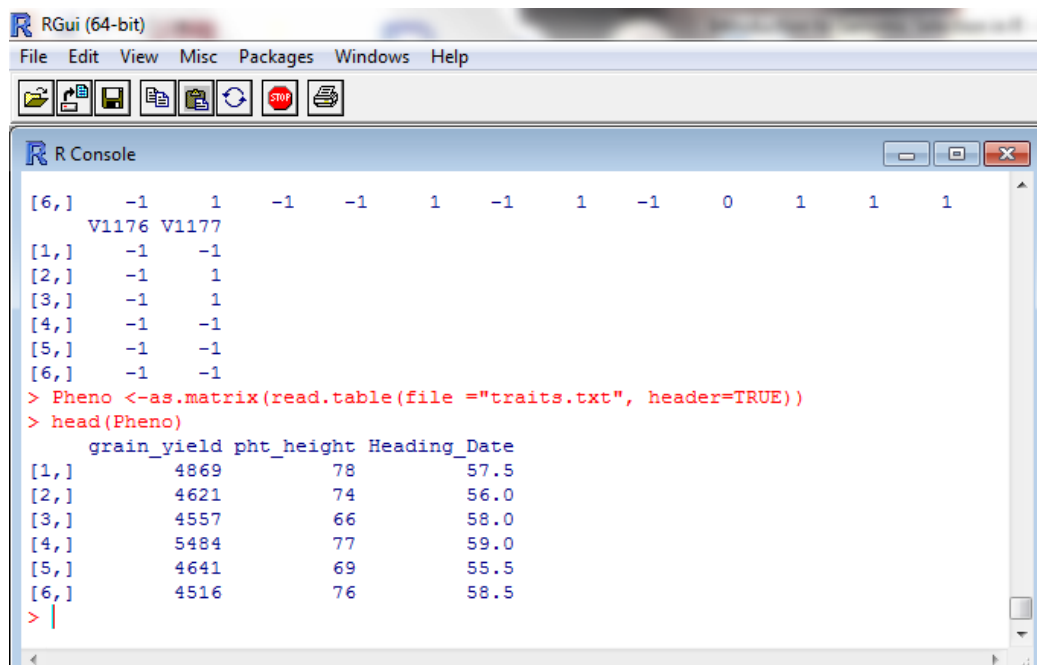


```
RGui (64-bit)
File Edit View Misc Packages Windows Help

R Console
> head(Markers)
      V1 V2 V3 V4 V5 V6 V7 V8 V9 V10 V11 V12 V13 V14 V15 V16 V17 V18 V19 V20 V21
[1,] -1  1  1 -1 -1  1 -1  1  1  1 -1  1  1  1 -1 -1 -1  1 -1  1  1
[2,] -1  1  1 -1 -1  1 -1  1  1  1  1  1  1  1  1 -1 -1  1 -1  1  1
[3,] -1  1  1 -1 -1  1 -1 -1 -1  1 -1  1  1  1 -1 -1 -1  1 -1  1  1
[4,] -1  1  1 -1 -1  1  1  1  1  1 -1  1  1  1 -1 -1 -1  1 -1  1  1
[5,]  1  1  1  1 -1  1 -1 -1 -1  1  1  1  1  1  1 -1 -1  1 -1  1  1
[6,] -1  1  1 -1 -1  1 -1 -1 -1  1 -1  1  1  1 -1 -1 -1  1 -1  1  1
      V22 V23 V24 V25 V26 V27 V28 V29 V30 V31 V32 V33 V34 V35 V36 V37 V38 V39
[1,] -1  1 -1  1 -1 -1 -1 -1  1  1 -1 -1  1  1 -1  1 -1  1
[2,] -1  1 -1  1 -1 -1 -1  1  1  1 -1 -1 -1  1 -1 -1 -1  1
[3,] -1  1 -1  1  1 -1 -1 -1  1  1 -1 -1 -1 -1 -1 -1 -1  1
[4,] -1  1 -1  1 -1 -1  1  1  1 -1 -1  1 -1 -1  1 -1  1  1
[5,] -1  1 -1  1  1 -1  1  1  1 -1 -1 -1  1 -1  1 -1  1  1
[6,] -1  1 -1  1 -1 -1  1  1  1 -1 -1  1  1 -1 -1 -1  1  1
      V40 V41 V42 V43 V44 V45 V46 V47 V48 V49 V50 V51 V52 V53 V54 V55 V56 V57
[1,] -1  1  1 -1 -1 -1 -1  1  1  1  1 -1 -1 -1 -1  1  1  1
[2,] -1  1  1 -1 -1 -1 -1  1  1  1  1 -1 -1 -1 -1 -1  1  1
```

Load the Sample Files

- Load the phenotype file and use the head command to see the first five lines
 - Header=T since phenotype files have column names
- Markers and phenotypes must be in matrix format

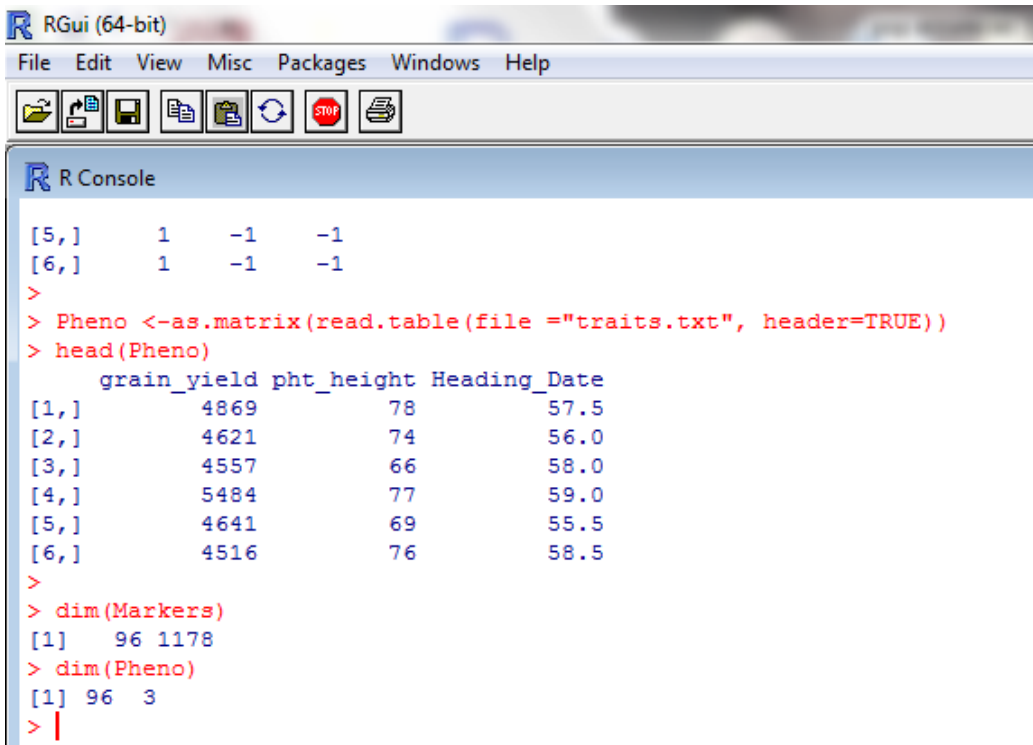


```
RGui (64-bit)
File Edit View Misc Packages Windows Help

R Console
[6,] -1 1 -1 -1 1 -1 1 -1 0 1 1 1
      V1176 V1177
[1,] -1 -1
[2,] -1 1
[3,] -1 1
[4,] -1 -1
[5,] -1 -1
[6,] -1 -1
> Pheno <- as.matrix(read.table(file = "traits.txt", header = TRUE))
> head(Pheno)
      grain_yield pht_height Heading_Date
[1,]      4869      78      57.5
[2,]      4621      74      56.0
[3,]      4557      66      58.0
[4,]      5484      77      59.0
[5,]      4641      69      55.5
[6,]      4516      76      58.5
> |
```

Load the Sample Files

- Determine the size of the matrices
- `dim()` command gives the number of rows and columns
- 96 observations and 1178 markers, 3 traits



The screenshot shows the RGui (64-bit) window. The menu bar includes File, Edit, View, Misc, Packages, Windows, and Help. The toolbar contains icons for file operations and execution. The R Console displays the following code and output:

```
> [5,] 1 -1 -1
> [6,] 1 -1 -1
> 
> Pheno <-as.matrix(read.table(file ="traits.txt", header=TRUE))
> head(Pheno)
      grain_yield pht_height Heading_Date
[1,]        4869         78         57.5
[2,]        4621         74         56.0
[3,]        4557         66         58.0
[4,]        5484         77         59.0
[5,]        4641         69         55.5
[6,]        4516         76         58.5
> 
> dim(Markers)
[1]  96 1178
> dim(Pheno)
[1] 96  3
> |
```

Learning Objectives

- Download the package and load the sample files
- **Impute missing markers using A.mat()**
- Define the training and validation populations
- Run mixed.solve() and determine accuracy of predictions

Impute Missing Markers

- rrBLUP mixed.solve() does not allow for missing markers
- Imputed value is the population mean for that marker
- Useful for SNP data since level of missing data is low
 - In the sample files 164 markers are missing out of 1178 (0.14%)
- A.mat also calculates the additive relationship matrix

Impute Missing Markers

- max.missing-maximum proportion of missing data
 - If 50% of markers are missing data then markers are not imputed
- impute method- imputes the mean of the markers
- return.imputed-prints out the imputed results if set to TRUE

```
> #what if markers are NA?  
> #impute with A.mat  
> impute=A.mat(Markers,max.missing=0.5,impute.method="mean",return.imputed=T)  
> |
```

Impute Missing Markers

- `>impute=A.mat(Markers,max.missing=0.5,impute.method="mean",return.imputed=T)`
- `> Markers_impute=impute$imputed`
 - Rename imputed marker matrix as `Markers_impute`
- `impute$imputed`-returns the imputed marker matrix
- `impute$A`-returns the additive relationship matrix

Impute Missing Markers

- `>impute$imputed`

Imputed
marker value

RGui (64-bit)

File Edit View Misc Packages Windows Help

R Console

	V137	V138	V139	V140	V141	V142	V143	V144	V145	V146	V147	V148	V149	V150
[1,]	-1	1	1	1	1	-1	-1	1	-1	0.8043478	1	1	1	1
[2,]	-1	1	1	1	1	-1	-1	-1	-1	1.0000000	1	1	1	1
[3,]	-1	1	1	1	1	-1	-1	1	-1	1.0000000	1	1	1	1
[4,]	1	-1	1	1	1	-1	-1	1	-1	1.0000000	1	1	1	1
[5,]	-1	1	1	1	1	-1	-1	-1	-1	1.0000000	1	-1	1	1
[6,]	-1	1	1	1	1	-1	-1	-1	-1	1.0000000	1	1	1	1

	V151	V152	V153	V154	V155	V156	V157	V158	V159	V160	V161	V162	V163	V164	V165
[1,]	1	1	-1	1	-1	1	1	-1	1	1	-1	-1	1	1	-1
[2,]	1	1	-1	1	-1	-1	1	-1	1	1	-1	-1	1	1	-1
[3,]	1	1	-1	1	-1	1	1	-1	1	1	-1	-1	1	1	-1
[4,]	1	1	-1	1	-1	1	1	-1	1	1	-1	-1	1	1	-1
[5,]	1	1	-1	1	-1	1	1	-1	1	1	-1	-1	1	1	1
[6,]	1	1	-1	1	-1	-1	1	-1	1	1	-1	-1	1	1	1

	V166	V167	V168	V169	V170	V171	V172	V173	V174	V175	V176	V177	V178	V179	V180
[1,]	1	1	1	NA	1	1	1	-1	-1	1	1	1	-1	-1	1
[2,]	1	1	1	NA	1	1	1	-1	-1	1	1	1	-1	-1	1
[3,]	1	1	1	NA	1	1	1	-1	1	1	1	1	-1	-1	1
[4,]	1	1	1	NA	1	1	1	-1	-1	1	1	1	-1	-1	1

Marker
value
left NA if
more
than
50%
missing
data

Impute Missing Markers

- Remove markers that had more than 50% missing data
 - NA values are not allowed in mixed.solve
 - Two markers in the SNP file must be removed
 - Column 169 and 562
 - New dimensions show 2 less columns
- Use Markers_impute2 as marker matrix for estimating marker effects

```
> Markers_impute2=Markers_impute[,-c(169,562)]  
> dim(Markers_impute)  
[1] 96 1178  
> dim(Markers_impute2)  
[1] 96 1176  
> |
```

Learning Objectives

- Download the package and load the sample files
- Impute missing markers using `A.mat()`
- **Define the training and validation populations**
- Run `mixed.solve()` and determine accuracy of predictions

Training and Validation Populations

- Training population-genotyped and phenotyped
- Validation population-phenotype values estimated based on marker effects calculated from training population
- Code is set that 60% of the total population is the training population
 - 40% validation population

Training and Validation Populations

- 58 (60% of total population of 96) random numbers sampled to determine which individuals are in the training population
- Individuals are the row numbers for the phenotypes and marker matrices
- Sampled numbers will be different every time the code is run and will affect the correlation accuracy

```
> train= as.matrix(sample(1:96, 58))
> head(train)
      [,1]
[1,]   52
[2,]   82
[3,]   50
[4,]   14
[5,]    7
[6,]   80
> |
```

Training and Validation Populations

- Validation population is 40% of the total population
- `setdiff()` command determines the numbers that are not in the training population and will be part of the validation population

```
> test<-setdiff(1:96,train)
> test
[1]  6 12 18 19 22 23 26 27 28 29 33 34 36 40 41 43 47 48 53 54 55 56 57 58 59 62 66
[28] 68 71 75 77 79 83 84 86 90 91 96
> |
```


Training and Validation Populations

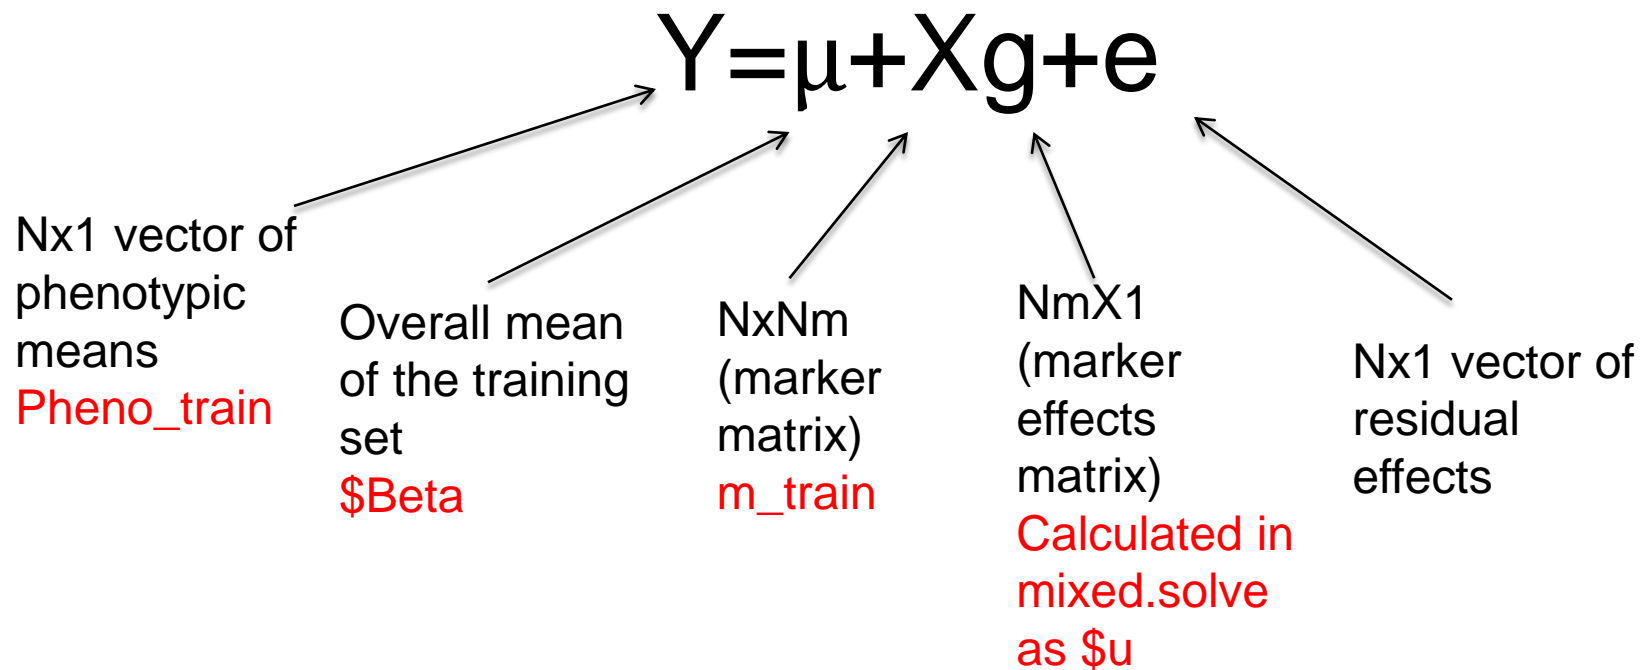
- Pheno_train and m_train are the phenotype and marker matrices for the values in the training population
- Pheno_valid and m_valid will be the validation populations

```
> Pheno_train=Pheno[train,]  
> m_train=Markers_impute2[train,]  
> Pheno_valid=Pheno[test,]  
> m_valid=Markers_impute2[test,]  
> |
```

Learning Objectives

- Download the package and load the sample files
- Impute missing markers using `A.mat()`
- Define the training and validation populations
- **Run `mixed.solve()` and determine accuracy of predictions**

Run mixed.solve



Run mixed.solve

Yield is the first column of the pheno_train matrix

```
> yield=(Pheno_train[,1])  
> yield_answer<-mixed.solve(yield, Z=m_train, K=NULL, SE = FALSE, return.Hinv=FALSE)  
> |
```

Vector of observations

Design matrix of random effects (Markers)

K matrix is the identity matrix

Standard errors are not calculated

Run mixed.solve

- Yield_answer\$u is the output of the marker effects
- head(e) shows the marker effects for the first five markers

```
> yield_answer<-mixed.solve(yield, Z=m_train, K=NULL, SE = FALSE, return.Hinv=FALSE)
> YLD = yield_answer$u
> e = as.matrix(YLD)
> head(e)
      [,1]
V1  1.1380597
V2 -0.1141220
V3  0.4970927
V4  2.2986051
V5  0.3579770
V6 -0.1141220
> |
```

Run mixed.solve

- $m_valid * e$ = marker validation matrix times the marker effects
- $Pred_yield$ = predicted yield based on the marker effects of the training population with the grand mean added in

```
> pred_yield_valid = m_valid %*% e
> pred_yield=(pred_yield_valid[,1])+yield_answer$beta
> pred_yield
[1] 4745.698 4621.133 4742.935 4601.210 4671.582 4636.899 4552.350 4486.954
[9] 4589.440 4601.534 4508.288 4656.675 4462.313 4493.898 4668.741 4498.701
[17] 4708.654 4593.296 4441.527 4705.500 4597.538 4089.056 4177.749 4261.560
[25] 4107.757 4207.431 4454.215 4713.850 4740.123 4537.690 4585.838 4526.935
[33] 4570.133 4512.558 4613.167 4412.658 4747.170 4872.127 4774.157 4697.992
[41] 4640.538 4576.519 4707.957 4658.228 4772.145 4596.747 4371.145 4779.256
[49] 4427.464 4525.557 4305.716 4564.654 4450.188 4634.591 3989.726 4068.685
[57] 4043.495 3886.869
```

Determine Correlation Accuracy

- Correlation between the predicted yield values and the observed yield values
- Accuracy will change slightly each time due to different individuals sampled for the training and validation populations

```
> yield_valid = Pheno_valid[,1]
> YLD_accuracy <-cor(pred_yield_valid, yield_valid, use="complete" )
> YLD_accuracy
      [,1]
[1,] 0.2521498
> |
```

Determine Correlation Accuracy

- Plant Height

```
> PHT_HT=(Pheno_train[,2])
> PHT_HT_answer<-mixed.solve(PHT_HT, Z=m_train, K=NULL, SE = FALSE, return.Hinv=FALSE)
> PHT_HT = PHT_HT_answer$u
> e = as.matrix(PHT_HT)
> pred_PHT_HT_valid = m_valid %*% e
> PHT_HT_valid = Pheno_valid[,2]
> PHT_HT_accuracy <-cor(pred_PHT_HT_valid, PHT_HT_valid, use="complete" )
> PHT_HT_accuracy
      [,1]
[1,] 0.4055428
> |
```


Determine Correlation Accuracy

- Heading Date

```
> HD_DATE=(Pheno_train[,3])
> HD_DATE_answer<-mixed.solve(HD_DATE, Z=m_train, SE = FALSE, return.Hinv=FALSE)
> HD_DATE = HD_DATE_answer$u
> e = as.matrix(HD_DATE)
> pred_HD_DATE_valid = m_valid %*% e
> HD_DATE_valid = Pheno_valid[,3]
> HD_DATE_accuracy <-cor(pred_HD_DATE_valid, HD_DATE_valid, use="complete" )
> HD_DATE_accuracy
      [,1]
[1,] 0.5205029
> |
```

Determine Correlation Accuracy

- Correlation accuracy with 500 iterations


```
>
> ##### cross validation for many cycles for yield only
> traits=1
> cycles=500
> accuracy = matrix(nrow=cycles, ncol=traits)
> for(r in 1:cycles)
+ {
+ train= as.matrix(sample(1:96, 38))
+ test<-setdiff(1:96,train)
+ Pheno_train=Pheno[train,]
+ m_train=Markers_impute2[train,]
+ Pheno_valid=Pheno[test,]
+ m_valid=Markers_impute2[test,]
+
+ yield=(Pheno_train[,1])
+ yield_answer<-mixed.solve(yield, Z=m_train, K=NULL, SE = FALSE, return.Hinv=F$
+ YLD = yield_answer$u
+ e = as.matrix(YLD)
+ pred_yield_valid = m_valid %*% e
+ pred_yield=(pred_yield_valid[,1])+yield_answer$beta
+ pred_yield
+ yield_valid = Pheno_valid[,1]
+ accuracy[r,1] <-cor(pred_yield_valid, yield_valid, use="complete" )
+ }
> mean(accuracy)
[1] 0.2305713
>
>
> |
```

Determine Correlation Accuracy

- Correlation accuracy is different for each trait
- Values will be different every time it is run since different lines will be included in the training or validation sets
- Accuracy is affected by training size, validation size, number of markers and heritability

Determine Correlation Accuracy

- Effects of training population size on accuracy



Common Errors

- Headers incorrectly input

```
> Pheno <-as.matrix(read.table(file ="traits.txt", header=F))
> head(Pheno)
      V1      V2      V3
[1,] "grain_yield" "pht_height" "Heading_Date"
[2,] "4869"        "78"        "57.5"
[3,] "4621"        "74"        "56"
[4,] "4557"        "66"        "58"
[5,] "5484"        "77"        "59"
[6,] "4641"        "69"        "55.5"
> train= as.matrix(sample(1:96, 38))
> test<-setdiff(1:96,train)
> Pheno_train=Pheno[train,]
> m_train=Markers_impute2[train,]
> Pheno_valid=Pheno[test,]
> m_valid=Markers_impute2[test,]
> yield=(Pheno_train[,1])
> yield_answer<-mixed.solve(yield, Z=m_train, K=NULL, SE = FALSE, return.Hinv=F$
Error in crossprod(x, y) :
  requires numeric/complex matrix/vector arguments
```

Common Errors

- NA Markers

```
> train= as.matrix(sample(1:96, 38))
> test<-setdiff(1:96,train)
> Pheno_train=Pheno[train,]
> m_train=Markers[train,]
> Pheno_valid=Pheno[test,]
> m_valid=Markers[test,]
> yield=(Pheno_train[,1])
> yield_answer<-mixed.solve(yield, Z=m_train, K=NULL, SE = FALSE, return.Hinv=FALSE)
Error in eigen(Hb, symmetric = TRUE) : infinite or missing values in 'x'
```

Common Errors

- Incorrect matrix dimensions
 - Removed one individual from phenotype matrix

```
> #####
> #define the training and test populations
> #training-60% validation-40%
> train= as.matrix(sample(1:96, 38))
> test<-setdiff(1:96,train)
> Pheno_train=Pheno[train,]
> m_train=Markers_impute2[train,]
> Pheno_valid=Pheno[test,]
Error: subscript out of bounds
> m_valid=Markers_impute2[test,]
>
> #####
> yield=(Pheno_train[,1])
> yield_answer<-mixed.solve(yield, Z=m_train, K=NULL, SE = FALSE, return.Hinv=F$
> YLD = yield_answer$u
> e = as.matrix(YLD)
> pred_yield_valid = m_valid %*% e
> pred_yield=(pred_yield_valid[,1])+yield_answer$beta
> yield_valid = Pheno_valid[,1]
Error: object 'Pheno_valid' not found
> YLD_accuracy <-cor(pred_yield_valid, yield_valid, use="complete" )
Error in is.data.frame(y) : object 'yield_valid' not found
> YLD_accuracy
Error: object 'YLD_accuracy' not found
> |
```

Common Errors

- Read in values as characters instead of numeric
 - Quotes around values

```
> Pheno <-as.matrix(read.table(file ="traits.txt", header=F))  
> head(Pheno)
```

	V1	V2	V3
[1,]	"grain_yield"	"pht_height"	"Heading_Date"
[2,]	"4869"	"78"	"57.5"
[3,]	"4621"	"74"	"56"
[4,]	"4557"	"66"	"58"
[5,]	"5484"	"77"	"59"

Resources

- rrBLUP reference manual
 - <http://cran.r-project.org/web/packages/rrBLUP/rrBLUP.pdf>
- rrBLUP vignettes
 - <http://cran.r-project.org/web/packages/rrBLUP/vignettes/vignette.pdf>
- Endelman, J.B. 2011. Ridge regression and other kernels for genomic selection with R package rrBLUP. Plant Genome 4:250-255. doi: 10.3835/plantgenome2011.08.0024

Acknowledgements

Rex Bernardo

Emily Combs

Lian Lian

Chris Schaefer

Lisa-Marie Krchov

rrBLUP

Jeff Endelman

Funding

Monsanto Company

USDA SoICAP

SoICAP

David Francis

Shawn Yarnes

John McQueen

Dataset

TCAP Hordeum's

Toolbox



MONSANTO



Questions?